

A CALIBRATED GALACTICUS GALAXY POPULATION MODEL FOR ROMAN GRISM SURVEY SIMULATIONS

ANDREW ROBERTSON* AND ANDREW BENSON

The Observatories of the Carnegie Institution for Science, 813 Santa Barbara Street, Pasadena, 91101, CA, USA
 Version August 13, 2026

ABSTRACT

Forward modelling the *Roman* Grism Survey requires mock galaxy catalogues that are simultaneously realistic in stellar mass, star formation, size, metallicity, black-hole demographics, and emission-line luminosity. We present a calibrated GALACTICUS semi-analytic model designed for this purpose. Because direct likelihood exploration with GALACTICUS is computationally expensive, we construct Gaussian-process emulators of binned observable predictions from a Sobol training campaign, compressing each observable with principal-component analysis where appropriate. We validate the emulators using held-out GALACTICUS runs, calibrate the model to a suite of low- and intermediate-redshift galaxy observables, and then extend the calibration to include observed H α luminosity functions. The final calibrated model reproduces the calibration data at the level required for Roman mock-catalogue generation. [TODO: Rewrite abstract once final claims and Roman quantities are fixed.]

1. INTRODUCTION

The Nancy Grace Roman Space Telescope will carry out a wide-field near-infrared grism survey over cosmological volumes, providing galaxy redshifts for large-scale structure measurements and a large sample of emission-line galaxies for studies of galaxy evolution (Spergel et al. 2015; Wang et al. 2022). These measurements will extend galaxy redshift surveys to large samples at redshifts where the expansion history and growth of structure are less tightly constrained by existing galaxy surveys. This is especially timely given recent DESI baryon-acoustic-oscillation (BAO) measurements. DESI first-year BAO data are consistent with Λ CDM on their own, but when combined with CMB and Type Ia supernova data they show a mild preference for time evolution of the dark energy equation of state (DESI Collaboration et al. 2025). This is not yet conclusive evidence for evolving dark energy, but it does sharpen the case for independent measurements of the distance-redshift relation and growth history. Roman is valuable in this context because it will probe these quantities at high redshift, with a different instrument, galaxy selection, and set of observational systematics from existing ground-based spectroscopic surveys.

Roman will obtain these redshifts in a very different way from a targeted spectroscopic survey such as DESI (DESI Collaboration et al. 2016). In DESI, robotic fibre positioners place fibre-optic cables on pre-selected galaxies and quasars, so that the spectrograph receives light only from those chosen locations on the sky. The Roman grism survey will instead be *slitless*: it will disperse the light from the full field of view, and will therefore obtain spectra for essentially every source that is bright enough to contribute detectable flux. This is a powerful feature of the survey, because it avoids a hard preselection of targets and can in principle provide spectra for all suffi-

ciently bright sources in the field. It also makes the data more complicated to interpret. Because the spectra are dispersed images of the galaxies, spatial and spectral information are mixed together, spectra from neighbouring objects can overlap, and the detectability of an emission line depends not only on its total flux but also on the size and morphology of the galaxy.

The redshifts used for the Roman grism survey will be obtained primarily from strong rest-frame optical emission lines, especially H α and [O III], with [O II] also relevant at higher redshift. For Roman mock catalogues, the line luminosities must be connected to the underlying galaxy population: stellar masses, star formation rates, gas metallicities, dust attenuation, galaxy sizes, and host halo environments. These correlations matter both for cosmological analyses, where the number density and bias of selected tracers are important, and for image-level simulations, where source sizes, line fluxes, continuum emission, and blending determine whether a redshift can be recovered.

There are several possible ways to build such mocks. Empirical models such as the UNIVERSEMACHINE (Behroozi et al. 2019) provide a powerful way to connect galaxies to dark matter halos, and newer empirical frameworks such as DIFFSKY extend this approach to star-formation histories and spectral-energy-distribution modelling. These models are extremely useful, but they are designed to be flexible descriptions of the galaxy-halo connection rather than direct calculations of the baryonic processes that regulate galaxy formation. Semi-analytic models offer a complementary route. In a semi-analytic model, galaxies are evolved along dark matter halo merger trees using approximate, physically motivated prescriptions for baryonic processes (e.g. Baugh 2006; Somerville & Davé 2015). These models are far less expensive than hydrodynamical simulations, but still encode explicit calculations for gas cooling, star formation, feedback, metal enrichment, galaxy sizes, black-hole

*email: arobertson@carnegiescience.edu

growth, and nebular emission. For Roman mock catalogues, this is attractive because the same model can be asked to predict many correlated properties of each galaxy.

In this paper we calibrate a GALACTICUS galaxy-formation model intended both for Roman grism mock generation and for broader applications requiring physically motivated, self-consistent galaxy catalogues. We make use of a broad set of calibration data because they constrain different aspects of the galaxy formation model. We focus on low and intermediate-redshift calibration data, because these measurements are typically better understood than the corresponding high-redshift constraints. At the same time, we include $H\alpha$ luminosity functions at higher redshift, because emission-line galaxies are central to the Roman grism survey and because a model calibrated only to the lower-redshift galaxy population is not necessarily adequate for this application.

GALACTICUS is much less expensive than a hydrodynamical simulation, but it is still too expensive to run directly inside an MCMC analysis. In Robertson & Benson (2026), we tried a different acceleration strategy: rather than evolving a large ensemble of merger trees, and the galaxies within them, for every likelihood evaluation, we calibrated GALACTICUS using small sets of galaxies at a few fixed halo masses. That approach is most natural for quantities for which there are observational estimates of the galaxy distribution at fixed halo mass, with stellar mass being the clearest example (e.g. observationally-inferred stellar-to-halo mass relations from Leauthaud et al. 2012; Behroozi et al. 2019). The resulting calibration produced a good match to the observed stellar mass function. One might hope that matching stellar mass functions at several redshifts would also constrain the predicted star-formation rates to match observations, because the model must form the correct amount of stellar mass between those redshifts. In practice, this is not sufficient: the integrated amount of star formation can be approximately correct even if the instantaneous distribution of star-formation rates is not. The model in Robertson & Benson (2026) therefore matched stellar mass functions more successfully than $H\alpha$ luminosity functions.

For the Roman grism survey this distinction matters, because the redshift sample is built from emission-line detections. Here we therefore use full GALACTICUS runs to compute the observables that enter the likelihood, including emission-line luminosity functions, but replace repeated direct model evaluations with an emulator. We construct this emulator from a finite training set of GALACTICUS runs, test its accuracy with held-out cross-validation runs, use it to sample the posterior distribution of the model parameters, and then check representative posterior models with direct GALACTICUS evaluations.

The paper is organized as follows. Section 2 describes the GALACTICUS model, the varied parameters, and the training campaign. Section 3 describes the emulator construction and validation. Section 4 describes the likelihood and MCMC inference. Section 5 presents the calibration results, including diagnostic subset fits and the final joint calibration. Section 6 discusses the interpretation and limitations of the calibration. We summarize in Section 7.

2. MODEL AND EMULATOR TRAINING SET

2.1. The GALACTICUS Model

We use GALACTICUS, a semi-analytic model of galaxy formation (Benson 2012). For a given dark matter halo accretion and merger history, encoded in the so-called *merger tree* (Lacey & Cole 1993), GALACTICUS follows the formation and evolution of the galaxies associated with each halo and subhalo. The calculation includes parameterised prescriptions for gas cooling, star formation, stellar feedback, AGN feedback, galaxy sizes, galaxy mergers, black-hole growth, and metal enrichment.

The fixed modelling choices used here are mostly the same as in Robertson & Benson (2026). By fixed modelling choices, we mean choices such as the prescriptions for disk star formation, gas cooling, black hole accretion, and stellar and AGN feedback. Compared with Robertson & Benson (2026), we vary a substantially larger set of numerical parameters within this model, as described in Section 2.3. We therefore focus here on the fixed modelling choices that differ from Robertson & Benson (2026) and on the parameters varied during the calibration. The GALACTICUS code is open source, and the complete parameter files used in this work will be made publicly available alongside the paper.¹

The main fixed model choice that differs from Robertson & Benson (2026) is the treatment of galaxy sizes. In that work, galaxy component sizes were computed using an iterative equilibrium calculation, in which the baryons contribute to the gravitational potential used to determine component radii. In the larger-volume runs needed here, this typically produced a small subpopulation of galaxies with unrealistically small sizes, with scale lengths roughly 10–100 times smaller than typical galaxies of the same stellar mass. We therefore use a simpler size calculation, in which component radii are determined from their specific angular momentum in the dark matter halo potential, without allowing the baryons to modify that potential.²

2.2. Merger Trees And Halo Sampling

For the training runs, GALACTICUS is run on extended Press–Schechter merger trees. We evolve the trees down to $z = 0$, where they contain 17,385 distinct host halos over the mass range $10^{11} < M_{\text{halo}}/M_{\odot} < 3 \times 10^{14}$. Their masses follow a Sheth–Tormen halo mass function (Sheth & Tormen 1999), with $a = 0.791$, $p = 0.218$, and normalisation 0.302, the parameter values from Benson et al. (2019). Each of these $z = 0$ host halos contains a central galaxy and may also host satellite galaxies. The calculation follows the complete merger tree of each $z = 0$ halo, including all progenitor branches. The abundance of host halos over this mass range corresponds to an effective comoving volume of $V_{\text{eff}} \simeq 1.3 \times 10^6 \text{ Mpc}^3$, equivalent to a cubic box with side length $\simeq 109 \text{ Mpc}$. Because the calibrated model is intended to generate Roman mocks using the UNIT simulations (Chuang et al. 2019), we use a merger-tree mass resolution of $2 \times 10^{10} M_{\odot}$, similar to

¹ [TODO: Add repository/version and parameter-file URL or Zenodo DOI.]

² Specifically, Robertson & Benson (2026) used `galacticStructureSolverEquilibrium`, whereas here we use `galacticStructureSolverSimple`; see the [Galacticus documentation](#).

the minimum halo mass resolved in UNIT. For the same reason, the merger trees are rebinned onto the UNIT simulation output times, following the approach of [Benson et al. \(2012\)](#).

The use of extended Press–Schechter trees is mainly a practical choice. It gives direct control over both the number of simulated halos and their mass distribution, which is useful when the goal is to build smooth training data for mass functions, size relations, emission-line luminosity functions, and other binned summary statistics. In Section 6.2 we include a comparison with a calculation based on merger trees extracted from the UNIT N -body simulation ([Chuang et al. 2019](#)), to assess the size of the effect for selected observables. [TODO: Update this sentence once the final EPS-versus-UNIT comparison figure and observable set are fixed.]

2.3. Calibrated Galacticus Parameters

Our aim is to calibrate a model suitable for mock-catalogue generation, for which accurately reproducing a broad range of observed galaxy properties is more important than reserving observables as independent tests of the model. We therefore vary 20 GALACTICUS parameters and constrain them using several galaxy observables simultaneously.

This differs from the calibration strategy often adopted for cosmological hydrodynamical simulations such as EAGLE, IllustrisTNG, and COLIBRE (e.g. [Schaye et al. 2015](#); [Nelson et al. 2018](#); [Schaye et al. 2026](#)), where a relatively small set of observables is typically used for calibration, leaving other galaxy properties as stringent tests of the model. While that is a valuable objective, it is not the primary goal of this work. Our 20 free parameters control aspect of gas cooling, star formation, stellar feedback, black-hole seeding, accretion and feedback, reincorporation of ejected gas, the classification of major and minor mergers and the sizes of merger remnants, angular-momentum transfer during disk instabilities, and the stellar metal yield. The parameters and their adopted ranges are listed in Table 1.

2.4. Emulator Training Set

The emulator training set consists of $N_{\text{train}} = 1024$ GALACTICUS model evaluations, with parameter values drawn from a scrambled Sobol sequence ([Sobol’ 1967](#); [Owen 1997](#)).³ Compared with random sampling, this provides a more uniform coverage of parameter space, reducing the likelihood of large unsampled regions and generally improving emulator accuracy for a fixed number of model evaluations. We generate Sobol points in the unit hypercube and transform each coordinate through the inverse cumulative distribution function of the corresponding parameter prior. As a result, the training density follows the prior density specified in Table 1, while retaining the space-filling properties of the Sobol design in prior-quantile space. The choice of 1024 training evaluations represents a compromise between reducing emulator error across the 20-dimensional parameter space and keeping the computational cost of the GALACTICUS runs manageable. Section 3.3 uses the low-redshift stel-

lar mass function as a concrete example to illustrate how emulator accuracy changes with training-set size.

2.5. Emission-Line And Dust Post-Processing

To connect the calibrated galaxy population to Roman grism observables, we augment the standard GALACTICUS outputs with emission-line luminosities and dust attenuation. For each GALACTICUS training run, we compute emission-line luminosity functions in post-processing using the `emission_line_dust` products. The current H α calibration uses the Sobral H α luminosity functions in four redshift samples. [TODO: Describe exactly how intrinsic line luminosities are computed, how dust is applied, and which dust parameters are varied. Add relation to Galacticus-dust-modelling package if appropriate.]

3. EMULATING CALIBRATION OBSERVABLES

3.1. Set Of Calibration Observables

Each GALACTICUS evaluation produces a catalogue of galaxies at a set of output times. From these galaxies we compute binned observable vectors, such as stellar mass functions, luminosity functions, size–mass relations, and scaling relations. The emulator is trained to reproduce these vectors as a function of the varied GALACTICUS parameters.

The final calibration uses the observable set listed in Table 2. For number-density observables we denote the binned differential number density by Φ , with the argument indicating the quantity being binned, for example $\Phi(M_\star)$ for a stellar mass function and $\Phi(L_{\text{H}\alpha})$ for an H α luminosity function. The emulated-quantity column of the table gives the transformed quantity passed to the emulator and used in the likelihood.

For each model evaluation, we also estimate the uncertainty on the predicted observable vector. This uncertainty arises because each observable is constructed from a finite simulated galaxy population, with the relevant uncertainty model depending on the observable being analysed. We use this information when training the emulator, so that noisier training outputs are treated as less precise estimates of the underlying model prediction. [TODO: Should describe the noise models used in a table or appendix or similar (or point to Galacticus documentation or similar).]

3.1.1. Observable-Space Transformations

For observables that span a large dynamic range across GALACTICUS parameter space, we emulate logarithmic quantities. For number-density observables this means emulating $\log_{10} \Phi$ rather than Φ itself. At fixed stellar mass, star formation rate, or emission-line luminosity, changes in the GALACTICUS parameters can produce order-of-magnitude changes in the predicted number density. The logarithmic transformation therefore gives a smoother emulator target across parameter space and makes emulator errors more closely related to fractional differences in Φ , matching the space in which luminosity and mass functions are usually compared.

We use the same approach for other positive quantities with broad dynamic range. The size–mass relations are emulated as $\log_{10} R_{\text{eff}}$ and the black-hole scaling relation as $\log_{10} M_{\text{BH}}$. The gas-phase mass–metallicity relation

³ The design was generated with `scipy.stats.qmc.Sobol`, using `scramble=True`, `seed=42`, and `random_base2(m=10)`.

Table 1

Parameters varied in the emulator training campaign. Here $\mathcal{LN}_{[a,b]}(x_0, \sigma_{\ln})$ denotes a log-normal prior truncated to $[a, b]$, with median x_0 and logarithmic width $\sigma_{\ln}$, while $\mathcal{U}_{[a,b]}$ denotes a uniform prior. Units for dimensional parameters are shown in the parameter column, and priors are given for the numerical value in those units.

Parameter	Description	Prior
Star formation and stellar populations		
$\nu_{\text{SF,d}}/\text{Gyr}^{-1}$	Disk star-formation frequency normalization	$\mathcal{LN}_{[0.1,10]}(1, 2)$
$\epsilon_{*,\text{sph}}$	Spheroid star-formation efficiency per dynamical time	$\mathcal{LN}_{[10^{-4},1]}(0.01, 4)$
Y_Z	Stellar metal yield	$\mathcal{LN}_{[0.01,0.04]}(0.02, 0.5)$
Stellar feedback		
$V_{\text{out,d}}/\text{km s}^{-1}$	Disk stellar-feedback velocity scale	$\mathcal{LN}_{[25,300]}(150, 0.5)$
$\alpha_{\text{out,d}}$	Disk stellar-feedback velocity exponent	$\mathcal{U}_{[0,4]}$
$V_{\text{out,sph}}/\text{km s}^{-1}$	Spheroid stellar-feedback velocity scale	$\mathcal{LN}_{[10,150]}(50, 1)$
Gas cooling and recycling		
f_{cool}	Gas cooling-rate multiplier	$\mathcal{LN}_{[0.05,1.2]}(0.5, 1)$
$r_{\text{core}}/R_{\text{vir}}$	Hot-gas core radius in units of virial radius	$\mathcal{LN}_{[0.03,1]}(0.3, 1)$
γ_{reinc}	Ejected-gas reincorporation normalization	$\mathcal{LN}_{[1,25]}(5, 0.5)$
$\delta_{\text{reinc},1}$	Redshift exponent in the Henriques reincorporation model	$\mathcal{U}_{[-2,4]}$
$\delta_{\text{reinc},2}$	Virial-velocity exponent in the Henriques reincorporation model	$\mathcal{U}_{[-1,5]}$
Mergers and disk instabilities		
f_{major}	Mass-ratio threshold for major mergers	$\mathcal{LN}_{[0.1,0.5]}(0.25, 0.5)$
f_{orbit}	Orbital-energy factor for merger-remnant sizes	$\mathcal{LN}_{[0.25,4]}(1, 0.5)$
$f_{j,\text{sph}}^{\text{bar}}$	Angular-momentum fraction retained by the spheroid in bar instabilities	$\mathcal{LN}_{[0.05,0.95]}(0.2, 0.5)$
Supermassive black holes		
$M_{\text{BH,seed}}/M_{\odot}$	Black-hole seed mass	$\mathcal{LN}_{[100,100\,000]}(3000, 4)$
$\epsilon_{\text{BH,radio}}$	Radio-mode black-hole heating efficiency	$\mathcal{LN}_{[0.03,1]}(0.3, 0.5)$
$\epsilon_{\text{BH,wind}}$	Black-hole wind feedback efficiency	$\mathcal{LN}_{[0.00024,0.024]}(0.0024, 1)$
$\alpha_{\text{Bondi,sph}}$	Bondi accretion enhancement for spheroid gas	$\mathcal{LN}_{[0.05,100]}(5, 2)$
$\alpha_{\text{Bondi,hot}}$	Bondi accretion enhancement for hot-halo gas	$\mathcal{LN}_{[0.05,100]}(6, 2)$
$x_{\text{thin,max}}$	Upper transition value for $\dot{M}c^2/L_{\text{Edd}}$ for switched thin-disk/ADAF model	$\mathcal{LN}_{[0.03,30]}(3, 2)$

Table 2

Calibration observables used in this work. The final joint fit uses all listed observables. N_{bin} is the number of binned data-vector elements entering the emulator and likelihood for the corresponding row. The emulated quantity defines the space used for emulator training and posterior inference, and is the space in which validation RMSE values are computed.

Calibration observable	Reference data	Emulated quantity	Redshift	N_{bin}
Stellar mass function	Tomczak et al. (2014)	$\log_{10} \Phi(M_{\star})$	$0.20 < z < 0.50$	14
			$1.00 < z < 1.25$	11
Star-formation-rate function	Robotham et al. (2011)	$\log_{10} \dot{\Phi}(M_{\star})$	$0.013 < z < 0.1$	25
Size-mass relation (star-forming)	van der Wel et al. (2014)	$\log_{10} R_{\text{eff}}(M_{\star})$	$0 < z < 0.5$	4
Size-mass relation (quiescent)		$\log_{10} R_{\text{eff}}(M_{\star})$	$0 < z < 0.5$	5
Black-hole-velocity-dispersion relation	McConnell & Ma (2013)	$\log_{10} M_{\text{BH}}(\sigma_{\star})$	$z \approx 0$	5
Gas-phase mass-metallicity relation	Blanc et al. (2019)	$12 + \log_{10}(\text{O}/\text{H})$	$z \approx 0$	49
			$z = 0.40$	18
H α luminosity function	Sobral et al. (2013)	$\log_{10} \Phi(L_{\text{H}\alpha})$	$z = 0.84$	9
			$z = 1.47$	13
			$z = 2.23$	15

is already conventionally expressed as $12 + \log_{10}(\text{O}/\text{H})$, and is emulated in that form.

Some training-set predictions are undefined after these transformations. For example, a model may produce no galaxies in a stellar-mass-function bin, so that $\log_{10} \Phi$ is undefined, or no quiescent galaxies in a particular stellar-mass bin, so that their median half-light radius is undefined. We treat these cases differently depending on the observable. For number densities, we impose a small positive floor before taking the logarithm. For undefined binned quantities such as median sizes, we replace the missing value with a representative value derived from the valid training-set predictions and assign it a large training uncertainty, so that the assigned value carries very little weight when fitting the emulator. The de-

tailed floor and missing-value procedure is described in Appendix A.3.

3.1.2. Principal-Component Compression

A simple approach to emulating each observable would be to build a separate emulator for every bin: for example, one emulator for the number density of galaxies at $M_{\star} \simeq 10^{10} M_{\odot}$, another at $M_{\star} \simeq 2 \times 10^{10} M_{\odot}$, and so on. This is straightforward, but it treats neighbouring bins as independent targets even though changes in GALACTICUS parameters usually move the predictions coherently across an observable. For example, some change in feedback strength will typically shift the amplitude, slope, or knee of a stellar mass function, rather than perturbing each stellar-mass bin independently.

We therefore use principal-component analysis (PCA)

to compress each observable vector before emulation. This reduces the number of emulators that must be trained and evaluated, and can improve predictive accuracy by exploiting correlations in how the different bins of an observable vary across parameter space. Principal-component compression is commonly used when emulating expensive simulation outputs with many correlated bins, including recent applications to cosmological and galaxy-formation inference (e.g. [Ramachandra et al. 2026](#)). For a given observable, let $\boldsymbol{\theta}$ denote the model parameters on which its prediction depends (namely, the GALACTICUS parameters, as well as potentially the dust model parameters), and let

$$\mathbf{y}(\boldsymbol{\theta}) = (y_1(\boldsymbol{\theta}), \dots, y_{N_{\text{bin}}}(\boldsymbol{\theta})) \quad (1)$$

be the vector of model predictions in the bins used for that observable. For the stellar mass function, for example, $y_b(\boldsymbol{\theta}) = \log_{10} \Phi_b(\boldsymbol{\theta})$, where Φ_b is the predicted number density in the b th stellar-mass bin. For each GALACTICUS run in the training set we have one such vector, so the training data for this observable form an $N_{\text{train}} \times N_{\text{bin}}$ matrix; we use n to index the individual GALACTICUS runs in the training set.

We first standardise each bin using its mean, μ_b , and standard deviation, s_b , across the training-set runs. For training-set run n , the standardised value in bin b is

$$z_{nb} = \frac{y_b(\boldsymbol{\theta}_n) - \mu_b}{s_b}. \quad (2)$$

We then apply PCA to the resulting $N_{\text{train}} \times N_{\text{bin}}$ matrix with elements z_{nb} . PCA defines an orthonormal set of directions, $\mathbf{e}_k$, in observable-bin space, ordered by the variance of the training set projected along each direction (e.g. [Jolliffe & Cadima 2016](#)). For training model n , its coefficient along component k is

$$a_{nk} = \mathbf{e}_k \cdot \mathbf{z}_n = \sum_b e_{kb} z_{nb}. \quad (3)$$

We retain the smallest number of components, K , required to explain 99 per cent of the variance in the standardised training set. The quantities a_{nk} are then the training targets for the emulator associated with component k ; after training, the emulator predicts the corresponding coefficient, $a_k(\boldsymbol{\theta})$, at new parameter values. The predicted observable vector in standardised observable space is reconstructed as

$$\mathbf{z}(\boldsymbol{\theta}) \simeq \sum_{k=1}^K a_k(\boldsymbol{\theta}) \mathbf{e}_k. \quad (4)$$

Finally, we undo the bin-wise standardisation to recover the emulator target in each bin,

$$y_b(\boldsymbol{\theta}) \simeq \mu_b + s_b z_b(\boldsymbol{\theta}). \quad (5)$$

For comparison and diagnostic plots, we also train bin-by-bin emulators. The main calibration results use PCA emulators with a 99 per cent variance threshold; the cross-validation tests below show why this is a reasonable choice.

3.2. Gaussian-Process Emulators

Each retained PCA coefficient is modelled with a Gaussian process as a function of the input parameters. The

same basic GP choices are used for the emulators compared below, so that the validation tests mainly probe the effects of training-set size and PCA compression. For the 20 GALACTICUS parameters, the GP inputs are not the physical parameter values themselves, but the corresponding prior quantiles. That is, for GALACTICUS parameter θ_i with prior cumulative distribution function $P_i(<\theta_i)$, we train the emulator using

$$u_i = P_i(<\theta_i), \quad (6)$$

so that each input parameter is mapped onto the unit interval. This is the space in which the Sobol design was constructed, and is therefore the space in which the training data are approximately space filling. It also avoids asking the GP to learn across raw parameters with very different units, ranges, and prior shapes; a fixed distance in u_i corresponds to the same amount of prior probability for each parameter. The H α luminosity-function emulator has five additional inputs: the dust post-processing parameters described in Section 2.5. Unlike the GALACTICUS parameters, these dust parameters are supplied to the GP as their model-parameter values, not as prior quantiles.

After the PCA transformation, each retained coefficient is standardised once more before fitting its GP, by subtracting its training-set mean and dividing by its training-set standard deviation. This is separate from the bin-wise standardisation applied before PCA, and is undone before reconstructing the observable vector. Let $\mathbf{u}$ denote the complete input vector supplied to a given GP after the input transformations just described. The kernel specifies the covariance between the values of the function at two points in this input space. We use a Matérn kernel with $\nu = 2.5$, multiplied by a constant amplitude, and allow a separate length scale for each input coordinate. For two input vectors, $\mathbf{u}$ and $\mathbf{u}'$, the kernel is

$$k(\mathbf{u}, \mathbf{u}') = A \left(1 + \sqrt{5}r + \frac{5r^2}{3} \right) \exp(-\sqrt{5}r), \quad (7)$$

where

$$r^2 = \sum_i \left(\frac{u_i - u'_i}{\ell_i} \right)^2. \quad (8)$$

Here A is the covariance amplitude and ℓ_i is the length scale associated with the i th emulator input. Small values of ℓ_i allow the emulator prediction to vary rapidly as that input coordinate is changed, while large values correspond to smoother variation. For each PCA coefficient, the kernel hyperparameters are therefore the overall covariance amplitude and one length scale for each emulator input. There are 20 length scales for the emulators of the non-H α calibration observables and 25 for the H α luminosity-function emulators, which also include the five dust parameters. We choose these hyperparameters by maximising the GP log marginal likelihood ([Rasmussen & Williams 2006](#)). In practice this maximisation is carried out with the L-BFGS-B algorithm, subject to the bounds on the kernel amplitude and length scales. Because the log marginal likelihood can have multiple local optima, the final emulators use five optimisations in total: one from the default initial kernel hyperparameters and four additional starting points randomly

drawn within the allowed hyperparameter bounds. We keep the solution with the largest log marginal likelihood. The final emulators do not include an additional fitted white-noise kernel; the estimated heteroscedastic training-output noise enters through the diagonal noise term described below.

3.2.1. Emulator Training Uncertainties

Where available, we include the estimated uncertainty in the training outputs by adding a diagonal noise term to the GP covariance matrix (e.g. [Rasmussen & Williams 2006](#)). In the notation used by `scikit-learn`, this diagonal variance contribution is the `alpha` parameter; here we refer to its square root as the training-output uncertainty. For emulator training, we represent the finite-sampling uncertainty of each binned observable by a per-bin variance in the transformed observable space. This corresponds to a diagonal training-noise approximation for the emulator: each training model and observable bin can have its own uncertainty, but inter-bin correlations are not propagated through the GP training. For example, the stellar mass function per-bin variances are estimated from a binomial finite-sampling model rather than from simple Poisson errors. This approximation is tested empirically by the held-out validation below.

For a PCA emulator, the per-bin training uncertainties must be propagated to the retained PCA coefficients. Let y_{nb} be the transformed observable for training model n in bin b , with variance σ_{nb}^2 . Using the bin-wise standardisation defined above,

$$z_{nb} = \frac{y_{nb} - \mu_b}{s_b}, \quad (9)$$

the corresponding variance is

$$\text{Var}(z_{nb}) = \frac{\sigma_{nb}^2}{s_b^2}. \quad (10)$$

The PCA coefficient for component k is

$$a_{nk} = \sum_b e_{kb} z_{nb}, \quad (11)$$

where e_{kb} is the b th element of the k th PCA component. Under the diagonal training-noise approximation, we assign this coefficient the variance

$$\text{Var}(a_{nk}) \simeq \sum_b e_{kb}^2 \frac{\sigma_{nb}^2}{s_b^2}. \quad (12)$$

Finally, because the scalar GP for coefficient k is fitted to the standardised coefficient

$$\tilde{a}_{nk} = \frac{a_{nk} - \bar{a}_k}{q_k}, \quad (13)$$

where $\bar{a}_k$ and q_k are the training-set mean and standard deviation of that coefficient, the variance supplied as `alpha` for that GP is $\text{Var}(a_{nk})/q_k^2$. For a bin-by-bin emulator, the analogous operation is simply to divide the per-bin transformed-space variance by the variance used to standardise the scalar GP target for that bin.

The distinction between emulator uncertainty and model uncertainty is important for the calibration. A GP prediction has both a mean and a predictive variance. The predictive variance is small near well-measured,

densely sampled regions of the training set, and grows where the training data are sparse, noisy, or require extrapolation. When this uncertainty is included in the likelihood, an imprecise emulator should therefore weaken the constraints on the GALACTICUS parameters rather than force the MCMC towards a spuriously precise answer. This is only true if the emulator uncertainty is reasonably calibrated, which is why the held-out validation tests in Section 3.2.2 are a necessary part of the inference rather than only a technical check.

3.2.2. Held-Out Emulator Validation

We validate the emulator using *five-fold cross-validation*. The GALACTICUS training runs are divided into five equally sized subsets. For each fold, the emulator is trained on four subsets and evaluated on the remaining subset, so that every GALACTICUS evaluation is predicted once by an emulator that was not trained using that model. Unless otherwise stated, all validation results in this section are based on the combined predictions from these five folds.

This has two purposes. First, it checks that the emulator can reproduce direct GALACTICUS calculations at new parameter values. Second, it lets us assess how choices made in constructing the emulator affect its predictive accuracy, including both the size of the training set and the representation of the observable vectors. Section 3.3 illustrates this validation procedure for the low-redshift stellar mass function.

To quantify this comparison, we consider individual bins of the observable vector. For a given bin, we compare the emulator prediction, m_i , with the direct GALACTICUS value, g_i , for each validation model i . The root-mean-square error,

$$\text{RMSE} = \left[\frac{1}{N_{\text{val}}} \sum_i (m_i - g_i)^2 \right]^{1/2}, \quad (14)$$

measures the typical vertical scatter about the one-to-one relation between emulator and direct GALACTICUS predictions. We also use the coefficient of determination,

$$R^2 = 1 - \frac{\sum_i (m_i - g_i)^2}{\sum_i (g_i - \bar{g})^2}, \quad (15)$$

where $\bar{g}$ is the mean direct GALACTICUS value over all validation models. This compares the emulator error with the intrinsic variation of the GALACTICUS predictions in that bin: $R^2 = 1$ corresponds to perfect predictions, while $R^2 = 0$ corresponds to doing no better than predicting the mean GALACTICUS value for all parameter choices. For the stellar mass functions shown below, both m_i and g_i are values of $\log_{10} \Phi$ in a fixed stellar-mass bin, so the RMSE is measured in dex.

3.3. Low-Redshift Stellar Mass Function Example

Figure 1 illustrates the held-out validation procedure for this observable. The grey curves show the range of GALACTICUS predictions across the training set, while the coloured points show several validation models. The corresponding coloured curves show the emulator predictions for the same parameter values.

Figure 2 shows this bin-level comparison for three representative stellar-mass bins from Figure 1. The plotted




Figure 1. Example of held-out validation for the low- z stellar mass function. Thin grey curves show the range of GALACTICUS predictions across the training set. Black points show the observed stellar mass function from Tomczak et al. (2014). Coloured points show direct GALACTICUS predictions for example held-out parameter choices. The corresponding coloured curves show emulator predictions for the same parameter choices, where in each case the emulator was trained without using the corresponding GALACTICUS evaluation. The shaded regions show the emulator uncertainty. The five example held-out models are selected only for visualisation: for each held-out model we compute the mean direct value of $\log_{10} \Phi$ across the plotted stellar-mass bins, rank the held-out models by this quantity, and choose examples at approximately evenly spaced quantiles of the ranking. The validation metrics use the full held-out sample. Vertical dotted lines mark the stellar-mass bins used for the bin-level validation metrics in Figures 2 and 3.

points should lie close to the one-to-one line if the emulator is accurate. The RMSE values measure the absolute scatter about this line, while the R^2 values measure how small that scatter is relative to the variation in the direct GALACTICUS predictions.

Figure 3 shows how these diagnostics depend on training-set size and emulator construction for several representative stellar-mass bins. We compare emulators trained on the first 64, 128, 256, 512, and 1024 points of the Sobol sequence. Because the first 2^m points of a Sobol sequence are themselves designed to provide good coverage of parameter space, these subsets provide a natural sequence of training sets for assessing the effect of increasing the training-set size. Emulator accuracy improves substantially as the training set grows from 64 to 1024 models, although the rate of improvement decreases for the larger training sets. This suggests that substantially reducing the remaining emulator error would require a much denser sampling of the 20-dimensional parameter space, with a correspondingly much larger training campaign. We therefore adopt the 1024-run training set as a practical compromise between emulator accuracy and computational cost.

Figure 3 also compares different levels of PCA compression. We retain 99 per cent of the PCA variance in the fiducial emulator, which keeps the number of GP models manageable while preserving the coherent response of neighbouring bins. We use this example to build intuition for the validation procedure; the final check on the emulator at the calibrated parameter values

is the direct GALACTICUS validation described in Section 5.3.

4. POSTERIOR INFERENCE OF MODEL PARAMETERS

4.1. MCMC Fits Used In This Work

We use two kinds of MCMC fit. First, we run a small set of diagnostic fits to subsets of the calibration data to illustrate how different observables constrain different combinations of GALACTICUS parameters. The example shown in Section 5.1 uses stellar mass functions, size-mass relations, and the star-formation-rate function, both separately and jointly. Second, we run the final joint calibration using the full set of calibration observables summarized in Table 2, including the Sobral $H\alpha$ luminosity functions. The final emulators are trained on the 1024-run Sobol campaign described in Section 2.4; no posterior-focused GALACTICUS evaluations are added to the emulator training set. The final joint fit samples the 20 GALACTICUS parameters in Table 1 together with the five dust parameters described in Appendix B.

4.2. The Likelihood

The MCMC likelihood is evaluated in the same transformed observable space in which the emulators are trained. For stellar mass functions, the star formation rate function, and the $H\alpha$ luminosity functions, this is $\log_{10}$ number density. For the size-mass relations and black-hole relation, the fitted quantities are already log-




Figure 2. Predicted versus held-out values of the low- z stellar mass function at the three stellar-mass bins highlighted in Figure 1. The x -axis values are the direct GALACTICUS predictions, while the y -axis shows the emulator predictions (with their uncertainty) at the same parameter values. The dashed lines show equality. The RMSE and R^2 values quoted in each panel summarise the accuracy of these held-out predictions, using the definitions given in Section 3.2.2. For clarity, the figure shows a random 25 per cent of the held-out comparisons.




Figure 3. Held-out emulator accuracy for the $z \simeq 0$ stellar mass function as a function of the number of GALACTICUS training runs. The top row shows the coefficient of determination, R^2 , and the bottom row shows the root-mean-square error in $\log_{10} \Phi$. The columns correspond to the representative stellar-mass bins highlighted in Figure 1. Different curves compare PCA-GP emulators with different PCA variance thresholds to a bin-by-bin GP emulator. The improvement from smaller training subsets to the full 1024-run training set motivates the larger final campaign used for the paper calibration, while the similarity of the PCA-GP variants supports the 99 per cent PCA variance threshold used in the fiducial emulator.

arithmetic quantities, $\log_{10} R_{\text{eff}}$ and $\log_{10} M_{\text{BH}}$, respectively. For the mass-metallicity relation, the fitted quantity is $12 + \log_{10}(\text{O}/\text{H})$.

For each observable bin i , we assume a Gaussian likelihood in this transformed space,

$$s_i^2(\theta) = \sigma_{\text{data},i}^2 + \sigma_{\text{emu},i}^2(\theta), \quad (16)$$

with

$$\ln \mathcal{L} = -\frac{1}{2} \sum_i \left[\frac{(d_i - m_i(\theta))^2}{s_i^2(\theta)} + \ln(2\pi s_i^2(\theta)) \right]. \quad (17)$$

Here d_i is the target value, $m_i(\theta)$ is the emulator mean prediction, $\sigma_{\text{data},i}$ is the observational uncertainty, and $\sigma_{\text{emu},i}(\theta)$ is the GP predictive uncertainty. The data and

emulator uncertainties are therefore added in quadrature in the transformed observable space.

Some of the GALACTICUS analysis outputs store full covariance matrices for the finite-volume uncertainty in the training predictions, and some also store full target covariance matrices. The final MCMC likelihood uses only the diagonal elements of these matrices, and therefore neglects correlations between bins and between different observable blocks. The observational uncertainties for logarithmic number densities are obtained by propagating the diagonal linear-space covariance into log space. A minimum uncertainty floor of 10^{-3} is imposed in the transformed observable space. The final calibration combines 113 non-H α observable bins with 55 H α luminosity-function bins. The H α likelihood is evalu-

ated in $\log_{10}[\Phi/(\text{Mpc}^{-3} \text{ dex}^{-1})]$, using the Sobral et al. target uncertainties propagated into this log-space quantity. Emulator uncertainties are included for all observable blocks in the final calibration.

For number-density observables whose target uncertainties are supplied in linear space, we convert the diagonal uncertainty to log space using first-order propagation,

$$\sigma_{\log_{10} \Phi} \simeq \frac{\sigma_{\Phi}}{\Phi \ln 10}. \quad (18)$$

For the Sobral et al. $\text{H}\alpha$ luminosity functions we instead use the tabulated uncertainties in $\log_{10} \Phi$ directly.

This gives a simple Gaussian likelihood in the transformed observable space. The approximation is most appropriate when the fractional uncertainties are moderate; sparsely populated number-density bins have intrinsically asymmetric counting uncertainties, which are only approximately represented by a symmetric Gaussian error bar in $\log_{10} \Phi$.

4.3. MCMC Sampling And Convergence

We sample the posterior distribution using Markov chain Monte Carlo (MCMC), implemented with the affine-invariant ensemble sampler `emcee` (Foreman-Mackey et al. 2013). Each MCMC uses 128 walkers, initialized in a small region around a point expected to lie close to the maximum of the posterior density. This point is determined from a preliminary `emcee` run with 128 walkers and 75,000 steps. During an initial burn-in phase, the walkers rapidly disperse from this small region and explore the posterior distribution. Initializing the walkers in this way reduces the likelihood of walkers becoming trapped around low-probability local maxima.

For the production runs, we use 250,000 steps per walker, discarding the initial 25,000 steps as burn-in. We assess convergence using the integrated autocorrelation time, τ (Sokal 1997), estimated separately for each sampled parameter. For the diagnostic fits used to illustrate parameter degeneracies (see Fig. 4), the post-burn-in chains extend for at least 25 times the largest measured autocorrelation time. For our final calibration, which jointly fits all calibration observables, the corresponding ratio is 73. We therefore consider the chains sufficiently long for the posterior inferences presented here.

The emulator predictions are used in place of direct GALACTICUS evaluations. This dramatically speeds up inference, because calling the emulator takes a few milliseconds, while a complete GALACTICUS evaluation takes hundreds of CPU hours.

5. CALIBRATION RESULTS

5.1. Degeneracy Breaking By Fitting Multiple Observables

To build intuition for the way in which calibrating to multiple observables helps break degeneracies between different GALACTICUS parameters, we first calibrate a model to three observable families:

- stellar mass functions from ZFOURGE/CANDELS (Tomczak et al. 2014) for redshifts $0.2 < z < 0.5$ and $1 < z < 1.25$,
- the low- z star-formation-rate function from GAMA (Robotham et al. 2011),

- the low- z size-mass relations for star-forming and quiescent galaxies from 3D-HST+CANDELS (van der Wel et al. 2014).

We fit these observables separately and jointly. Comparing the resulting posteriors shows where individual observables leave degeneracies and where the combined data prefer a smaller region of parameter space.

5.2. Final Joint Calibration

We then sample the posterior using all calibration observables, combining the emulator for the non- $\text{H}\alpha$ observables with the $\text{H}\alpha$ luminosity-function emulator. The $\text{H}\alpha$ luminosity-function emulator is trained from emission-line luminosity functions computed by dust post-processing the GALACTICUS training runs. For the fiducial configuration, we train on the four Sobral redshift samples, retain PCA components explaining 99 per cent of the variance, use GP training uncertainties, and optimise GP hyperparameters as described in Section 3. [TODO: Confirm final dust-draw choice and the exact emission-line luminosity-function training-row accounting.]

5.3. Direct Galacticus Validation At The MAP Point

Figure 7 compares the emulator prediction at the maximum-a-posteriori point to a direct GALACTICUS run at the same parameter values for the non- $\text{H}\alpha$ calibration observables. Figure 8 shows the corresponding comparison for the Sobral $\text{H}\alpha$ luminosity functions. These direct evaluations test the full model at the inferred parameter values rather than only the emulator.

5.4. Interpretation Of The Calibrated Model

The calibration demonstrates that a single GALACTICUS parameter set can reproduce a broad set of galaxy-population observables while also matching the observed $\text{H}\alpha$ luminosity functions used to anchor Roman grism mocks. [TODO: Use this section to interpret which observables are well matched, where the calibrated model compromises, and which parameter shifts are physically meaningful.]

6. DISCUSSION

6.1. What Emulator-Assisted Calibration Enables

The emulator makes posterior inference cheap enough to repeat for different combinations of calibration data. This allows us to compare subset fits, identify which observables drive particular parameter constraints, and locate tensions that would be difficult to diagnose with direct GALACTICUS evaluations alone. In future applications, the same framework can be extended by adding new observable blocks and retraining or augmenting the relevant emulators.

6.2. Comparison With N-Body Merger Trees

The calibration runs use extended Press-Schechter trees because they allow controlled, repeatable sampling of the halo mass function. This choice is not meant to imply that EPS trees capture every aspect of halo assembly in an N -body simulation. [TODO: Add comparison between the fiducial EPS-tree calculation and an otherwise similar calculation using N -body merger trees. The




Figure 4. An example of combining multiple observables to break degeneracies in the emulator calibration. Left: posterior constraints on a subset of six model parameters from fits to the stellar mass functions alone, the star formation rate function alone, galaxy sizes alone, and their combination. The grey dashed curves in the 1D marginal panels show the adopted priors. Right: shaded regions show the central 68 per cent intervals of the emulator-predicted mean relations obtained by drawing parameters from each posterior, compared with the observational data. These bands reflect posterior uncertainty in the calibrated model parameters, and do not include observational noise or emulator uncertainty. Faint grey curves show the predictions from the 1024 training runs, illustrating the range spanned by the prior over model parameters. The combined calibration retains good agreement with all fitted observables while reducing the parameter degeneracies present in the single-observable fits.

Figure 5. Posterior constraints from individual diagnostic calibrations and their joint fit. [TODO: Use polished version of the GetDist overlay. Consider putting the full 20-parameter version in an appendix and a reduced set in the main text.]

Figure 6. Summary of the final calibrated model across the full set of calibration observables. [TODO: This may be redundant with Figures 7 and 8; decide once figure drafts exist.]

Figure 7. Posterior predictive comparison for the non- $H\alpha$ calibration observables. [TODO: Desired styling: data points; emulator MAP or posterior band; direct Galacticus MAP validation for the final model.]

Figure 8. $H\alpha$ luminosity-function comparison for the final joint calibration. [TODO: Use final MAP/predictive plot with emulator uncertainty and direct Galacticus MAP validation.]

useful figure is probably not a parameter-inference figure, but a direct observable comparison showing whether the EPS choice changes the calibrated observables at a level that matters for this paper.]

6.3. Limitations

Several limitations should be kept explicit. First, the emission-line luminosity functions used here are part of the calibration data, so Roman number-density predictions are not independent tests of the $H\alpha$ luminosity function. Second, the emulator is trained over a finite parameter volume, and posterior inferences near prior boundaries or poorly sampled regions should be treated cautiously. Third, the likelihood assumes a particular

treatment of observational covariance, emulator uncertainty, and possible systematics between heterogeneous datasets. Fourth, the final mock catalogues will depend not only on the calibrated GALACTICUS model but also on lightcone construction, survey selection, spectral contamination, detector effects, and image simulation.

6.4. Future Work

The next step is to populate a $4 \times 4 \text{ deg}^2$ lightcone with mock galaxies drawn from the calibrated model and pass the resulting catalogue through Roman grism image simulations. Those image-level mocks will provide a public downstream data product and will support analysis development within the Roman Project Infrastructure

Team.

Inference being cheap allows one to try fitting to different subsets of calibration data, and see where tensions arise, as done in Bose & Deason (2026).

[TODO: Give CPU time for generating training data (mean run times was 15 hours on 16 Caltech CPUs), say that we could afford to use more, but (as seen for the SMF in the example section) emulator error decreases fairly slowly with increasing Ntrain. Maybe state that earlier iterations of our emulator used 512 training samples, with emulator uncertainties that were not that much smaller. State that in the future one could “target” promising regions of the parameter space, for additional training runs, with some references to this being done in the literature. Give some sort of intuition for the fact that if you could decrease the prior range by a factor of 2 in each of 20 dimensions, this would be a factor of a million increase in the density of training points, while running a million times more models would be infeasible.]

7. CONCLUSIONS

1. We constructed a 20-parameter emulator training campaign of 1024 GALACTICUS evaluations for galaxy-population and emission-line mock calibration, using a scrambled Sobol design and extended Press–Schechter trees with controlled sampling of the $z = 0$ halo mass function.
2. We built PCA-GP emulators of binned observable vectors, using prior-quantile inputs, standardised observable outputs, PCA compression, and Gaussian-process regression of the retained coefficients.
3. Held-out GALACTICUS evaluations and training-size tests validate the emulator for posterior inference over the calibration observables.
4. The calibration constrains the model using stellar mass functions, the star-formation-rate function, size–mass relations, the black-hole–halo-mass relation, the mass–metallicity relation, and H α luminosity functions.
5. Adding H α luminosity functions produces a final calibrated model suitable for Roman grism mock-catalogue generation, with direct GALACTICUS MAP runs confirming the emulator-based solution.
6. The calibrated posterior provides the starting point for future lightcone and image simulations of the Roman grism survey.

[TODO: Rewrite conclusions after figures are final; keep to 4–6 claims.]

ACKNOWLEDGEMENTS

[TODO: Add funding, Roman PIT acknowledgement language, and HPC acknowledgement.]

SOFTWARE

This work made use of [TODO: Astropy, emcee, GetDist, h5py, Matplotlib, NumPy, pandas, scikit-learn, SciPy, and any Galacticus-specific packages. Add citations.]

REFERENCES

- Baugh, C. M. 2006, Reports on Progress in Physics, 69, 3101, doi: [10.1088/0034-4885/69/12/R02](#)
- Behroozi, P., Wechsler, R. H., Hearin, A. P., & Conroy, C. 2019, MNRAS, 488, 3143, doi: [10.1093/mnras/stz1182](#)
- Benson, A. J. 2012, New A, 17, 175, doi: [10.1016/j.newast.2011.07.004](#)
- Benson, A. J., Borgani, S., De Lucia, G., Boylan-Kolchin, M., & Monaco, P. 2012, MNRAS, 419, 3590, doi: [10.1111/j.1365-2966.2011.20002.x](#)
- Benson, A. J., Ludlow, A., & Cole, S. 2019, MNRAS, 485, 5010, doi: [10.1093/mnras/stz695](#)
- Blanc, G. A., Lu, Y., Benson, A., Katsianis, A., & Barraza, M. 2019, ApJ, 877, 6, doi: [10.3847/1538-4357/ab16ec](#)
- Bose, S., & Deason, A. J. 2026, MNRAS, 546, stag198, doi: [10.1093/mnras/stag198](#)
- Chuang, C.-H., Yepes, G., Kitaura, F.-S., et al. 2019, MNRAS, 487, 48, doi: [10.1093/mnras/stz1233](#)
- DESI Collaboration, Adame, A. G., et al. 2025, Journal of Cosmology and Astroparticle Physics, 2025, 021, doi: [10.1088/1475-7516/2025/02/021](#)
- DESI Collaboration, Aghamousa, A., Aguilar, J., et al. 2016, arXiv e-prints, arXiv:1611.00036, doi: [10.48550/arXiv.1611.00036](#)
- Foreman-Mackey, D., Hogg, D. W., Lang, D., & Goodman, J. 2013, PASP, 125, 306, doi: [10.1086/670067](#)
- Jolliffe, I. T., & Cadima, J. 2016, Philosophical Transactions of the Royal Society of London Series A, 374, 20150202, doi: [10.1098/rsta.2015.0202](#)
- Lacey, C., & Cole, S. 1993, MNRAS, 262, 627, doi: [10.1093/mnras/262.3.627](#)
- Leauthaud, A., Tinker, J., Bundy, K., et al. 2012, ApJ, 744, 159, doi: [10.1088/0004-637X/744/2/159](#)
- McConnell, N. J., & Ma, C.-P. 2013, ApJ, 764, 184, doi: [10.1088/0004-637X/764/2/184](#)
- Nelson, D., Pillepich, A., Springel, V., et al. 2018, MNRAS, 475, 624, doi: [10.1093/mnras/stx3040](#)
- Owen, A. B. 1997, The Annals of Statistics, 25, 1541, doi: [10.1214/aos/1031594731](#)
- Ramachandra, N., Frontiere, N., Buehlmann, M., et al. 2026, arXiv e-prints, arXiv:2601.07306, doi: [10.48550/arXiv.2601.07306](#)
- Rasmussen, C. E., & Williams, C. K. I. 2006, Gaussian Processes for Machine Learning
- Robertson, A., & Benson, A. 2026, The Open Journal of Astrophysics, 9, 55, doi: [10.33232/001c.155306](#)
- Robotham, A. S. G., Norberg, P., Driver, S. P., et al. 2011, MNRAS, 416, 2640, doi: [10.1111/j.1365-2966.2011.19217.x](#)
- Schaye, J., Crain, R. A., Bower, R. G., et al. 2015, MNRAS, 446, 521, doi: [10.1093/mnras/stu2058](#)
- Schaye, J., Chaikin, E., Schaller, M., et al. 2026, MNRAS, 548, stag375, doi: [10.1093/mnras/stag375](#)
- Sheth, R. K., & Tormen, G. 1999, MNRAS, 308, 119, doi: [10.1046/j.1365-8711.1999.02692.x](#)
- Sobol', I. M. 1967, USSR Computational Mathematics and Mathematical Physics, 7, 86, doi: [10.1016/0041-5553\(67\)90144-9](#)
- Sobral, D., Smail, I., Best, P. N., et al. 2013, MNRAS, 428, 1128, doi: [10.1093/mnras/sts096](#)
- Sokal, A. 1997, Monte Carlo Methods in Statistical Mechanics: Foundations and New Algorithms, ed. C. DeWitt-Morette, P. Cartier, & A. Folacci (Boston, MA: Springer US), 131–192, doi: [10.1007/978-1-4899-0319-8_6](#)
- Somerville, R. S., & Davé, R. 2015, ARA&A, 53, 51, doi: [10.1146/annurev-astro-082812-140951](#)
- Spergel, D., Gehrels, N., Baltay, C., et al. 2015, arXiv e-prints, arXiv:1503.03757, doi: [10.48550/arXiv.1503.03757](#)
- Tomczak, A. R., Quadri, R. F., Tran, K.-V. H., et al. 2014, ApJ, 783, 85, doi: [10.1088/0004-637X/783/2/85](#)
- van der Wel, A., Franx, M., van Dokkum, P. G., et al. 2014, ApJ, 788, 28, doi: [10.1088/0004-637X/788/1/28](#)
- Wang, Y., Zhai, Z., Alavi, A., et al. 2022, ApJ, 928, 1, doi: [10.3847/1538-4357/ac4973](#)

Table 3

Five-fold held-out validation summary for the non-H α calibration-observable emulators trained from the full 1024-run campaign. Each row summarizes the PCA-GP emulator for one observable block. N_{bin} and N_{PC} correspond to the full emulator trained on all 1024 GALACTICUS runs. For the cross-validation calculation, the full emulator construction is repeated independently in each fold. The remaining entries are medians over the bins retained by the full emulator. RMSE is measured in the transformed observable space used by the emulator and likelihood. $f_{1\sigma}$ and $f_{2\sigma}$ give the fractions of held-out predictions whose direct GALACTICUS values lie within the corresponding emulator uncertainty interval. The H α LF emulator is not included because an analogous grouped-CV summary has not yet been generated for the production emission-line luminosity-function emulator.

Observable block	N_{bin}	N_{PC}	RMSE	R^2	RMSE/ σ_{emu}	$f_{1\sigma}$	$f_{2\sigma}$
Black-hole-velocity-dispersion relation	5	4	0.32	0.79	1.47	0.62	0.88
ZFOURGE stellar mass function, $0.20 < z < 0.50$	14	6	0.11	0.92	1.24	0.69	0.92
ZFOURGE stellar mass function, $1.00 < z < 1.25$	11	4	0.24	0.92	1.36	0.63	0.89
Blanc et al. gas-phase mass-metallicity relation	49	6	0.08	0.89	1.46	0.61	0.89
Robotham et al. star-formation-rate function	25	7	0.11	0.88	1.11	0.69	0.93
van der Wel et al. size-mass relation, star-forming	4	3	0.05	0.90	1.15	0.69	0.92
van der Wel et al. size-mass relation, quiescent	5	4	0.06	0.84	1.19	0.70	0.92

APPENDIX

A. OBSERVABLE AND EMULATOR DETAILS

A.1. *Observable Transformations And Compression*

[TODO: Add per-observable transformation rules and PCA component counts selected by the 99 per cent variance threshold.]

A.2. *Held-Out Emulator Validation Summary*

Table 3 summarizes the held-out tests for the non-H α observable emulators used in the calibration. If the quoted emulator uncertainties were exactly right, the RMSE divided by the emulator error would be close to one, and about 68 and 95 per cent of held-out predictions would fall within the quoted one- and two-sigma error bars. The table is close to this expectation, but the RMSE/error values are usually a little larger than one and the two-sigma coverage is slightly below 95 per cent. This means that the emulator error bars are somewhat too small in these tests. That level of mismatch is not surprising: the GP error bars do not include every approximation in the emulator construction, such as the finite choice of kernel, the finite PCA basis, and the approximate treatment of noisy training outputs. We therefore treat the held-out tests as a check on the scale of the emulator error, and use the direct GALACTICUS MAP validation in Section 5.3 as a complementary check at the parameter point used for the final calibration.

A.3. *Treatment Of Missing, Zero, Or Non-Positive Training Values*

[TODO: Document the exact handling of non-finite, zero, or non-positive training values, including any replacement values and inflated training uncertainties.]

A.4. *Covariance And Uncertainty Treatment*

[TODO: Document which observables have non-zero off-diagonal target covariance, which have non-zero off-diagonal finite-volume covariance in the GALACTICUS analysis outputs, and where only diagonal terms are used by the emulator and likelihood.]

B. EMISSION-LINE DUST MODEL DETAILS

[TODO: Describe the intrinsic emission-line luminosity calculation, the dust attenuation law, the five dust parameters, their priors, and the post-processing used for the Sobral H α luminosity functions.]

This paper was built using the Open Journal of Astrophysics L^AT_EX template. The OJA is a journal which provides fast and easy peer review for new papers in the `astro-ph` section of the arXiv, making the reviewing process simpler for authors and referees alike. Learn more at <http://astro.theoj.org>.